

In [86]:

```
from keras.utils import np_utils #import utilities from the keras
from keras.datasets import mnist #loading mnist dataset
import seaborn as sns #seaborn library for plotting
from keras.initializers import RandomNormal #to initialize weights we have an inbuilt keras
library and from that we are importing randomnormal initializer
```

In [87]:

```
%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [88]:

```
# the data, shuffled and split between train and test sets
#loading the mnist data
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

In [89]:

```
X_train.shape
```

Out[89]:

```
(60000, 28, 28)
```

In [90]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

```
Number of training examples : 60000 and each image is of shape (28, 28)
```

```
Number of training examples : 10000 and each image is of shape (28, 28)
```

In [91]:

```
# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [92]:

```
# after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

```
Number of training examples : 60000 and each image is of shape (784)
```

Number of training examples : 10000 and each image is of shape (784, 784, 3)

In [93]:

```
# An example data point
print(X_train[0])
```

[	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	3	18	18	18	126	136	175	26	166	255	
247	127	0	0	0	0	0	0	0	0	0	0	0	0	0	30	36	94	154	
170	253	253	253	253	253	225	172	253	242	195	64	0	0	0	0	0	0	0	
	0	0	0	0	0	49	238	253	253	253	253	253	253	253	253	251	93	82	
82	56	39	0	0	0	0	0	0	0	0	0	0	0	0	0	18	219	253	
253	253	253	253	198	182	247	241	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	80	156	107	253	253	205	11	0	43	154		
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	14	1	154	253	90	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	139	253	190	2	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	11	190	253	70	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	35	241	
225	160	108	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	81	240	253	253	119	25	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	45	186	253	253	150	27	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	16	93	252	253	187	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	249	253	249	64	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	46	130	183	253	
253	207	2	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	39	148	229	253	253	253	250	182	0	0	0	0	0	0	
	0	0	0	0															

The above matrix is a greyscale image 28\*28 pixel image with values ranging from 0 to 255 with 0 as white and as we move from 0 to 255 it turns from white to grey to black

In [94]:

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 

X_train = X_train/255
X_test = X_test/255
```

In [95]:

```
# example data point after normlizing
print(X_train[0])
```

[0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.

[illegible]



In [97]:

```
# https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.Each Layer has exactly one input tensor and one
output tensor
# you can create a Sequential model by passing a list of layer instances to the constructor:

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None,
activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT. X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument s
upported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions ar available ex: tanh, relu, softmax

from keras.models import Sequential #importing a sequential model
from keras.layers import Dense, Activation #importing dense and also activation to use activation
functions
```

In [98]:

```
# some model parameters

output_dim = 10 #since we have to class numbers from 0 to 9 we have output dimension as 10
input_dim = X_train.shape[1] #input dimension will the image i.e. 784(28*28 px image)
print(input_dim)

batch_size = 128 #used a batch size of 128
nb_epoch = 20 #number of epochs is 28
```

# ReLU + Adam Optimization + 2 Hidden Layers (without Dropout and Batch Normalization)

Hidden layer 1 consists of 256 outputs and Hidden Layer2 consists of 64 connections

In [99]:

```
model_relu = Sequential()
model_relu.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Model: "sequential\_16"

Layer (type)	Output Shape	Param #
dense_59 (Dense)	(None, 256)	200960
dense_60 (Dense)	(None, 64)	16448
dense_61 (Dense)	(None, 10)	650
Total params: 218,058		
Trainable params: 218,058		
Non-trainable params: 0		

None  
Train on 60000 samples, validate on 10000 samples  
Epoch 1/20  
60000/60000 [=====] - 3s 55us/step - loss: 0.2779 - accuracy: 0.9176 - val\_loss: 0.1286 - val\_accuracy: 0.9611  
Epoch 2/20  
60000/60000 [=====] - 3s 53us/step - loss: 0.1073 - accuracy: 0.9682 - val\_loss: 0.0953 - val\_accuracy: 0.9702  
Epoch 3/20  
60000/60000 [=====] - 3s 52us/step - loss: 0.0721 - accuracy: 0.9781 - val\_loss: 0.0790 - val\_accuracy: 0.9757  
Epoch 4/20  
60000/60000 [=====] - 3s 53us/step - loss: 0.0501 - accuracy: 0.9842 - val\_loss: 0.0791 - val\_accuracy: 0.9752  
Epoch 5/20  
60000/60000 [=====] - 3s 53us/step - loss: 0.0373 - accuracy: 0.9883 - val\_loss: 0.0819 - val\_accuracy: 0.9763  
Epoch 6/20  
60000/60000 [=====] - 3s 53us/step - loss: 0.0301 - accuracy: 0.9905 - val\_loss: 0.0742 - val\_accuracy: 0.9780  
Epoch 7/20  
60000/60000 [=====] - 3s 53us/step - loss: 0.0215 - accuracy: 0.9937 - val\_loss: 0.0788 - val\_accuracy: 0.9775  
Epoch 8/20  
60000/60000 [=====] - 3s 53us/step - loss: 0.0167 - accuracy: 0.9951 - val\_loss: 0.0702 - val\_accuracy: 0.9801  
Epoch 9/20  
60000/60000 [=====] - 3s 53us/step - loss: 0.0148 - accuracy: 0.9956 - val\_loss: 0.0791 - val\_accuracy: 0.9803  
Epoch 10/20  
60000/60000 [=====] - 3s 52us/step - loss: 0.0136 - accuracy: 0.9955 - val\_loss: 0.0770 - val\_accuracy: 0.9803  
Epoch 11/20  
60000/60000 [=====] - 3s 52us/step - loss: 0.0112 - accuracy: 0.9964 - val\_loss: 0.0793 - val\_accuracy: 0.9796  
Epoch 12/20  
60000/60000 [=====] - 3s 52us/step - loss: 0.0101 - accuracy: 0.9969 - val\_loss: 0.0943 - val\_accuracy: 0.9771  
Epoch 13/20

```
Epoch 13/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0114 - accuracy: 0.9962 - va
l_loss: 0.0854 - val_accuracy: 0.9793
Epoch 14/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0077 - accuracy: 0.9978 - va
l_loss: 0.0798 - val_accuracy: 0.9820
Epoch 15/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0038 - accuracy: 0.9990 - va
l_loss: 0.0814 - val_accuracy: 0.9826
Epoch 16/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0094 - accuracy: 0.9966 - va
l_loss: 0.0946 - val_accuracy: 0.9789
Epoch 17/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0083 - accuracy: 0.9971 - va
l_loss: 0.0990 - val_accuracy: 0.9793
Epoch 18/20
60000/60000 [=====] - 3s 54us/step - loss: 0.0066 - accuracy: 0.9977 - va
l_loss: 0.1108 - val_accuracy: 0.9767
Epoch 19/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0057 - accuracy: 0.9983 - va
l_loss: 0.0828 - val_accuracy: 0.9820
Epoch 20/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0047 - accuracy: 0.9987 - va
l_loss: 0.0900 - val_accuracy: 0.9808
```

In [100]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.08999803272318979  
Test accuracy: 0.9807999730110168



after epoch 5 there is not much change in the validation loss whereas our train loss keeps on decreasing around 14 epochs after there is not much change

In [101]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10,10))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## ReLU + Adam Optimization + 3 Hidden Layers (without Dropout and Batch Normalization)

Hidden layer 1 consists of 512 outputs and Hidden Layer 2 consists of 256 outputs. Hidden Layer



hidden layer 1 consists of 512 outputs and hidden Layer2 consists of 256 outputs, hidden Layer 3 consists of 128 outputs

In [102]:

```
# Multilayer perceptron

model_relu = Sequential() #says we want a sequential model
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()#prints the summary
```

Model: "sequential\_17"

Layer (type)	Output Shape	Param #
dense_62 (Dense)	(None, 512)	401920
dense_63 (Dense)	(None, 256)	131328
dense_64 (Dense)	(None, 128)	32896
dense_65 (Dense)	(None, 10)	1290
Total params: 567,434		
Trainable params: 567,434		
Non-trainable params: 0		

In [103]:

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 7s 121us/step - loss: 0.2373 - accuracy: 0.9292 - val_loss: 0.1105 - val_accuracy: 0.9657
Epoch 2/20
60000/60000 [=====] - 7s 118us/step - loss: 0.0871 - accuracy: 0.9735 - val_loss: 0.0991 - val_accuracy: 0.9703
Epoch 3/20
60000/60000 [=====] - 7s 119us/step - loss: 0.0554 - accuracy: 0.9830 - val_loss: 0.0782 - val_accuracy: 0.9770
Epoch 4/20
60000/60000 [=====] - 7s 118us/step - loss: 0.0397 - accuracy: 0.9877 - val_loss: 0.0740 - val_accuracy: 0.9783
Epoch 5/20
60000/60000 [=====] - 7s 118us/step - loss: 0.0319 - accuracy: 0.9891 - val_loss: 0.0766 - val_accuracy: 0.9777
Epoch 6/20
60000/60000 [=====] - 7s 118us/step - loss: 0.0251 - accuracy: 0.9920 - val_loss: 0.0751 - val_accuracy: 0.9794
Epoch 7/20
60000/60000 [=====] - 7s 119us/step - loss: 0.0227 - accuracy: 0.9925 - val_loss: 0.0886 - val_accuracy: 0.9780
Epoch 8/20
60000/60000 [=====] - 7s 118us/step - loss: 0.0198 - accuracy: 0.9936 - val_loss: 0.0940 - val_accuracy: 0.9756
Epoch 9/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0166 - accuracy: 0.9944 - val_loss: 0.0876 - val_accuracy: 0.9764
Epoch 10/20
60000/60000 [=====] - 7s 119us/step - loss: 0.0186 - accuracy: 0.9940 - val_loss: 0.0976 - val_accuracy: 0.9757
Epoch 11/20
60000/60000 [=====] - 7s 122us/step - loss: 0.0184 - accuracy: 0.9942 - v
```

```

al_loss: 0.0989 - val_accuracy: 0.9760
Epoch 12/20
60000/60000 [=====] - 7s 123us/step - loss: 0.0167 - accuracy: 0.9948 - v
al_loss: 0.0877 - val_accuracy: 0.9793
Epoch 13/20
60000/60000 [=====] - 7s 122us/step - loss: 0.0104 - accuracy: 0.9966 - v
al_loss: 0.0914 - val_accuracy: 0.9805
Epoch 14/20
60000/60000 [=====] - 7s 122us/step - loss: 0.0108 - accuracy: 0.9964 - v
al_loss: 0.0889 - val_accuracy: 0.9811
Epoch 15/20
60000/60000 [=====] - 7s 118us/step - loss: 0.0127 - accuracy: 0.9959 - v
al_loss: 0.0988 - val_accuracy: 0.9780
Epoch 16/20
60000/60000 [=====] - 7s 121us/step - loss: 0.0120 - accuracy: 0.9959 - v
al_loss: 0.1108 - val_accuracy: 0.9777
Epoch 17/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0103 - accuracy: 0.9969 - v
al_loss: 0.1004 - val_accuracy: 0.9808
Epoch 18/20
60000/60000 [=====] - 7s 119us/step - loss: 0.0137 - accuracy: 0.9957 - v
al_loss: 0.0970 - val_accuracy: 0.9785
Epoch 19/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0094 - accuracy: 0.9970 - v
al_loss: 0.0974 - val_accuracy: 0.9810
Epoch 20/20
60000/60000 [=====] - 7s 119us/step - loss: 0.0062 - accuracy: 0.9980 - v
al_loss: 0.1026 - val_accuracy: 0.9801

```

In [104]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.10256570401383279  
Test accuracy: 0.9800999760627747





after epoch 3 we can see jagged lines in our validation loss whereas our train loss keeps on decreasing around 15 epochs after there is not much change

In [105]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10,10))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



**ReLU + Adam Optimization + 5 Hidden Layers (without Dropout**

## and Batch Normalization)

**Hidden layer 1 consists of 512 outputs and Hidden Layer2 consists of 256 outputs, Hidden Layer 3 consists of 128 outputs ,Hidden Layer 4 consists of 64 outputs,Hidden Layer 5 consists of 32 outputs**

In [106]:

```
# Multilayer perceptron

model_relu = Sequential() #says we want a sequential model
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary() #prints the summary
```

Model: "sequential\_18"

Layer (type)	Output Shape	Param #
dense_66 (Dense)	(None, 512)	401920
dense_67 (Dense)	(None, 256)	131328
dense_68 (Dense)	(None, 128)	32896
dense_69 (Dense)	(None, 64)	8256
dense_70 (Dense)	(None, 32)	2080
dense_71 (Dense)	(None, 10)	330
Total params: 576,810		
Trainable params: 576,810		
Non-trainable params: 0		

In [107]:

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 130us/step - loss: 0.2782 - accuracy: 0.9158 - val\_loss: 0.1288 - val\_accuracy: 0.9611

Epoch 2/20

60000/60000 [=====] - 8s 127us/step - loss: 0.0949 - accuracy: 0.9704 - val\_loss: 0.1105 - val\_accuracy: 0.9667

Epoch 3/20

60000/60000 [=====] - 8s 127us/step - loss: 0.0642 - accuracy: 0.9791 - val\_loss: 0.0897 - val\_accuracy: 0.9722

Epoch 4/20

60000/60000 [=====] - 8s 125us/step - loss: 0.0463 - accuracy: 0.9853 - val\_loss: 0.0894 - val\_accuracy: 0.9732

Epoch 5/20

60000/60000 [=====] - 8s 126us/step - loss: 0.0377 - accuracy: 0.9880 - val\_loss: 0.0833 - val\_accuracy: 0.9776

Epoch 6/20

60000/60000 [=====] - 8s 126us/step - loss: 0.0296 - accuracy: 0.9906 - val\_loss: 0.0791 - val\_accuracy: 0.9789

```

Epoch 7/20
60000/60000 [=====] - 8s 127us/step - loss: 0.0283 - accuracy: 0.9907 - v
al_loss: 0.0828 - val_accuracy: 0.9765
Epoch 8/20
60000/60000 [=====] - 7s 125us/step - loss: 0.0226 - accuracy: 0.9925 - v
al_loss: 0.0909 - val_accuracy: 0.9756
Epoch 9/20
60000/60000 [=====] - 7s 125us/step - loss: 0.0209 - accuracy: 0.9933 - v
al_loss: 0.0809 - val_accuracy: 0.9797
Epoch 10/20
60000/60000 [=====] - 7s 124us/step - loss: 0.0199 - accuracy: 0.9936 - v
al_loss: 0.1038 - val_accuracy: 0.9758
Epoch 11/20
60000/60000 [=====] - 7s 123us/step - loss: 0.0193 - accuracy: 0.9938 - v
al_loss: 0.1011 - val_accuracy: 0.9755
Epoch 12/20
60000/60000 [=====] - 7s 123us/step - loss: 0.0141 - accuracy: 0.9953 - v
al_loss: 0.0958 - val_accuracy: 0.9781
Epoch 13/20
60000/60000 [=====] - 7s 124us/step - loss: 0.0159 - accuracy: 0.9948 - v
al_loss: 0.1009 - val_accuracy: 0.9780
Epoch 14/20
60000/60000 [=====] - 7s 124us/step - loss: 0.0198 - accuracy: 0.9940 - v
al_loss: 0.1104 - val_accuracy: 0.9759
Epoch 15/20
60000/60000 [=====] - 7s 124us/step - loss: 0.0123 - accuracy: 0.9959 - v
al_loss: 0.0900 - val_accuracy: 0.9829
Epoch 16/20
60000/60000 [=====] - 7s 123us/step - loss: 0.0124 - accuracy: 0.9961 - v
al_loss: 0.0914 - val_accuracy: 0.9806
Epoch 17/20
60000/60000 [=====] - 7s 123us/step - loss: 0.0147 - accuracy: 0.9952 - v
al_loss: 0.0968 - val_accuracy: 0.9795
Epoch 18/20
60000/60000 [=====] - 7s 124us/step - loss: 0.0105 - accuracy: 0.9968 - v
al_loss: 0.1007 - val_accuracy: 0.9789
Epoch 19/20
60000/60000 [=====] - 7s 125us/step - loss: 0.0083 - accuracy: 0.9974 - v
al_loss: 0.1163 - val_accuracy: 0.9785
Epoch 20/20
60000/60000 [=====] - 7s 121us/step - loss: 0.0119 - accuracy: 0.9965 - v
al_loss: 0.0762 - val_accuracy: 0.9823

```

In [108]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

```

Test score: 0.07623335870537294
Test accuracy: 0.9822999835014343

```



In [109]:

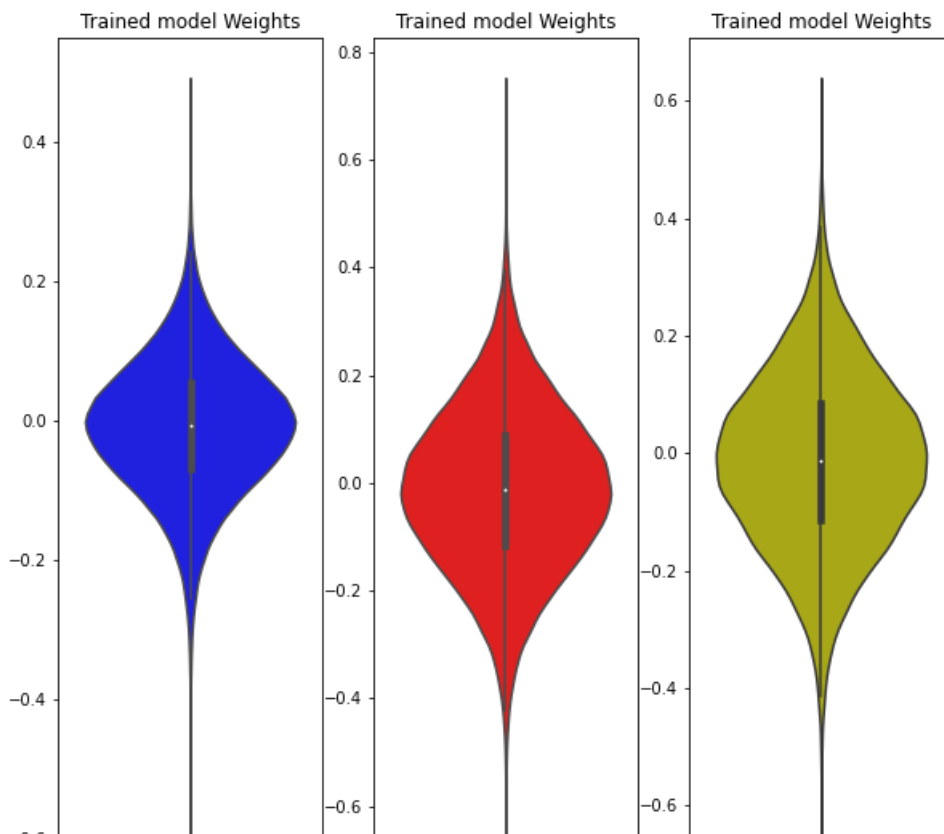
```
w_after = model_relu.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10,10))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





## ReLU + Adam Optimization + 2 Hidden Layers (Batch Normalization)

In [110]:

```
from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

Model: "sequential\_19"

Layer (type)	Output Shape	Param #
dense_72 (Dense)	(None, 256)	200960
batch_normalization_23 (Batch Normalization)	(None, 256)	1024
dense_73 (Dense)	(None, 64)	16448
batch_normalization_24 (Batch Normalization)	(None, 64)	256
dense_74 (Dense)	(None, 10)	650
Total params: 219,338		
Trainable params: 218,698		
Non-trainable params: 640		

In [111]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 4s 68us/step - loss: 0.2419 - accuracy: 0.9304 - val_loss: 0.1283 - val_accuracy: 0.9619
Epoch 2/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0953 - accuracy: 0.9719 - val_loss: 0.0932 - val_accuracy: 0.9709
Epoch 3/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0639 - accuracy: 0.9810 - val_loss: 0.0881 - val_accuracy: 0.9730
Epoch 4/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0484 - accuracy: 0.9849 - val_loss: 0.0840 - val_accuracy: 0.9735
Epoch 5/20
60000/60000 [=====] - 4s 61us/step - loss: 0.0363 - accuracy: 0.9890 - val_loss: 0.0958 - val_accuracy: 0.9703
Epoch 6/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0284 - accuracy: 0.9913 - val_loss: 0.0958 - val_accuracy: 0.9703
```

```

l_loss: 0.0855 - val_accuracy: 0.9730
Epoch 7/20
60000/60000 [=====] - 4s 61us/step - loss: 0.0219 - accuracy: 0.9933 - va
l_loss: 0.0829 - val_accuracy: 0.9753
Epoch 8/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0191 - accuracy: 0.9941 - va
l_loss: 0.0784 - val_accuracy: 0.9767
Epoch 9/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0181 - accuracy: 0.9943 - va
l_loss: 0.0885 - val_accuracy: 0.9748
Epoch 10/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0167 - accuracy: 0.9950 - va
l_loss: 0.0941 - val_accuracy: 0.9746
Epoch 11/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0140 - accuracy: 0.9956 - va
l_loss: 0.0875 - val_accuracy: 0.9748
Epoch 12/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0139 - accuracy: 0.9959 - va
l_loss: 0.0889 - val_accuracy: 0.9765
Epoch 13/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0120 - accuracy: 0.9962 - va
l_loss: 0.0796 - val_accuracy: 0.9796
Epoch 14/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0090 - accuracy: 0.9975 - va
l_loss: 0.0767 - val_accuracy: 0.9796
Epoch 15/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0073 - accuracy: 0.9979 - va
l_loss: 0.0940 - val_accuracy: 0.9764
Epoch 16/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0095 - accuracy: 0.9972 - va
l_loss: 0.0963 - val_accuracy: 0.9777
Epoch 17/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0095 - accuracy: 0.9973 - va
l_loss: 0.0828 - val_accuracy: 0.9787
Epoch 18/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0069 - accuracy: 0.9980 - va
l_loss: 0.0913 - val_accuracy: 0.9773
Epoch 19/20
60000/60000 [=====] - 4s 61us/step - loss: 0.0089 - accuracy: 0.9971 - va
l_loss: 0.1049 - val_accuracy: 0.9735
Epoch 20/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0072 - accuracy: 0.9975 - va
l_loss: 0.0935 - val_accuracy: 0.9767

```

In [112]:

```

score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)


```

```

Test score: 0.09346026394422224
Test accuracy: 0.9767000079154968

```





In [113]:


```
w_after = model_batch.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10,10))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





## ReLU + Adam Optimization + 3 Hidden Layers (Batch Normalization)

In [114]:

```
from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

Model: "sequential\_20"

Layer (type)	Output Shape	Param #
dense_75 (Dense)	(None, 512)	401920
batch_normalization_25 (Batch Normalization)	(None, 512)	2048
dense_76 (Dense)	(None, 256)	131328
batch_normalization_26 (Batch Normalization)	(None, 256)	1024
dense_77 (Dense)	(None, 128)	32896
batch_normalization_27 (Batch Normalization)	(None, 128)	512
dense_78 (Dense)	(None, 10)	1290
Total params: 571,018		
Trainable params: 569,226		
Non-trainable params: 1,792		

In [115]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 9s 148us/step - loss: 0.2114 - accuracy: 0.9378 - val_loss: 0.1119 - val_accuracy: 0.9654
Epoch 2/20
60000/60000 [=====] - 9s 142us/step - loss: 0.0754 - accuracy: 0.9776 - val_loss: 0.0913 - val_accuracy: 0.9715
Epoch 3/20
60000/60000 [=====] - 8s 141us/step - loss: 0.0473 - accuracy: 0.9857 - val_loss: 0.0888 - val_accuracy: 0.9721
```

```

Epoch 4/20
60000/60000 [=====] - 9s 143us/step - loss: 0.0349 - accuracy: 0.9884 - v
al_loss: 0.0789 - val_accuracy: 0.9750
Epoch 5/20
60000/60000 [=====] - 8s 141us/step - loss: 0.0280 - accuracy: 0.9912 - v
al_loss: 0.0849 - val_accuracy: 0.9740
Epoch 6/20
60000/60000 [=====] - 9s 142us/step - loss: 0.0206 - accuracy: 0.9934 - v
al_loss: 0.0852 - val_accuracy: 0.9756
Epoch 7/20
60000/60000 [=====] - 8s 141us/step - loss: 0.0176 - accuracy: 0.9942 - v
al_loss: 0.0805 - val_accuracy: 0.9767
Epoch 8/20
60000/60000 [=====] - 8s 141us/step - loss: 0.0164 - accuracy: 0.9945 - v
al_loss: 0.0709 - val_accuracy: 0.9791
Epoch 9/20
60000/60000 [=====] - 9s 142us/step - loss: 0.0153 - accuracy: 0.9948 - v
al_loss: 0.0816 - val_accuracy: 0.9769
Epoch 10/20
60000/60000 [=====] - 9s 144us/step - loss: 0.0138 - accuracy: 0.9955 - v
al_loss: 0.0832 - val_accuracy: 0.9765
Epoch 11/20
60000/60000 [=====] - 9s 148us/step - loss: 0.0114 - accuracy: 0.9961 - v
al_loss: 0.0807 - val_accuracy: 0.9767
Epoch 12/20
60000/60000 [=====] - 9s 151us/step - loss: 0.0107 - accuracy: 0.9964 - v
al_loss: 0.0827 - val_accuracy: 0.9774
Epoch 13/20
60000/60000 [=====] - 9s 145us/step - loss: 0.0117 - accuracy: 0.9961 - v
al_loss: 0.0898 - val_accuracy: 0.9793
Epoch 14/20
60000/60000 [=====] - 9s 144us/step - loss: 0.0105 - accuracy: 0.9968 - v
al_loss: 0.0833 - val_accuracy: 0.9792
Epoch 15/20
60000/60000 [=====] - 9s 144us/step - loss: 0.0077 - accuracy: 0.9974 - v
al_loss: 0.0857 - val_accuracy: 0.9790
Epoch 16/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0100 - accuracy: 0.9965 - v
al_loss: 0.0808 - val_accuracy: 0.9809
Epoch 17/20
60000/60000 [=====] - 9s 144us/step - loss: 0.0105 - accuracy: 0.9966 - v
al_loss: 0.0823 - val_accuracy: 0.9813
Epoch 18/20
60000/60000 [=====] - 9s 144us/step - loss: 0.0070 - accuracy: 0.9975 - v
al_loss: 0.0790 - val_accuracy: 0.9809
Epoch 19/20
60000/60000 [=====] - 9s 143us/step - loss: 0.0070 - accuracy: 0.9977 - v
al_loss: 0.0827 - val_accuracy: 0.9776
Epoch 20/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0067 - accuracy: 0.9978 - v
al_loss: 0.0769 - val_accuracy: 0.9819

```

In [116]:

```

score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

```

```
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.07694103825540023  
Test accuracy: 0.9818999767303467



In [117]:


```
w_after = model_batch.get_weights()

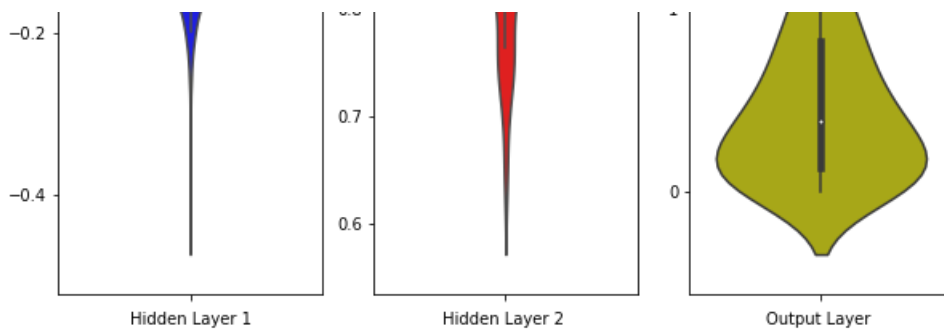
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10,10))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





## ReLU + Adam Optimization + 5 Hidden Layers (Batch Normalization)

In [118]:

```
from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

Model: "sequential\_21"

Layer (type)	Output Shape	Param #
dense_79 (Dense)	(None, 512)	401920
batch_normalization_28 (Batch Normalization)	(None, 512)	2048
dense_80 (Dense)	(None, 256)	131328
batch_normalization_29 (Batch Normalization)	(None, 256)	1024
dense_81 (Dense)	(None, 128)	32896
batch_normalization_30 (Batch Normalization)	(None, 128)	512
dense_82 (Dense)	(None, 64)	8256
batch_normalization_31 (Batch Normalization)	(None, 64)	256
dense_83 (Dense)	(None, 32)	2080
batch_normalization_32 (Batch Normalization)	(None, 32)	128
dense_84 (Dense)	(None, 10)	330

Total params: 580,778  
Trainable params: 578,794  
Non-trainable params: 1,984

---

In [119]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 10s 166us/step - loss: 0.2957 - accuracy: 0.9186 - val\_loss: 0.1249 - val\_accuracy: 0.9658

Epoch 2/20

60000/60000 [=====] - 9s 152us/step - loss: 0.0979 - accuracy: 0.9710 - val\_loss: 0.1022 - val\_accuracy: 0.9691

Epoch 3/20

60000/60000 [=====] - 9s 152us/step - loss: 0.0676 - accuracy: 0.9790 - val\_loss: 0.0939 - val\_accuracy: 0.9710

Epoch 4/20

60000/60000 [=====] - 9s 153us/step - loss: 0.0509 - accuracy: 0.9843 - val\_loss: 0.0959 - val\_accuracy: 0.9682

Epoch 5/20

60000/60000 [=====] - 9s 153us/step - loss: 0.0365 - accuracy: 0.9881 - val\_loss: 0.0879 - val\_accuracy: 0.9741

Epoch 6/20

60000/60000 [=====] - 9s 151us/step - loss: 0.0320 - accuracy: 0.9897 - val\_loss: 0.0871 - val\_accuracy: 0.9760

Epoch 7/20

60000/60000 [=====] - 9s 153us/step - loss: 0.0268 - accuracy: 0.9916 - val\_loss: 0.0845 - val\_accuracy: 0.9756

Epoch 8/20

60000/60000 [=====] - 9s 154us/step - loss: 0.0270 - accuracy: 0.9910 - val\_loss: 0.0773 - val\_accuracy: 0.9787

Epoch 9/20

60000/60000 [=====] - 9s 153us/step - loss: 0.0214 - accuracy: 0.9930 - val\_loss: 0.0816 - val\_accuracy: 0.9771

Epoch 10/20

60000/60000 [=====] - 9s 153us/step - loss: 0.0183 - accuracy: 0.9940 - val\_loss: 0.0951 - val\_accuracy: 0.9723

Epoch 11/20

60000/60000 [=====] - 9s 155us/step - loss: 0.0183 - accuracy: 0.9936 - val\_loss: 0.0888 - val\_accuracy: 0.9754

Epoch 12/20

60000/60000 [=====] - 9s 154us/step - loss: 0.0166 - accuracy: 0.9944 - val\_loss: 0.0880 - val\_accuracy: 0.9768

Epoch 13/20

60000/60000 [=====] - 9s 152us/step - loss: 0.0146 - accuracy: 0.9954 - val\_loss: 0.0904 - val\_accuracy: 0.9759

Epoch 14/20

60000/60000 [=====] - 9s 155us/step - loss: 0.0174 - accuracy: 0.9944 - val\_loss: 0.0770 - val\_accuracy: 0.9795

Epoch 15/20

60000/60000 [=====] - 9s 151us/step - loss: 0.0149 - accuracy: 0.9951 - val\_loss: 0.0869 - val\_accuracy: 0.9779

Epoch 16/20

60000/60000 [=====] - 9s 155us/step - loss: 0.0112 - accuracy: 0.9964 - val\_loss: 0.0912 - val\_accuracy: 0.9780

Epoch 17/20

60000/60000 [=====] - 10s 160us/step - loss: 0.0113 - accuracy: 0.9961 - val\_loss: 0.0812 - val\_accuracy: 0.9808

Epoch 18/20

60000/60000 [=====] - 9s 153us/step - loss: 0.0110 - accuracy: 0.9963 - val\_loss: 0.0908 - val\_accuracy: 0.9785

Epoch 19/20

60000/60000 [=====] - 9s 153us/step - loss: 0.0119 - accuracy: 0.9962 - val\_loss: 0.0916 - val\_accuracy: 0.9802

Epoch 20/20

60000/60000 [=====] - 9s 151us/step - loss: 0.0109 - accuracy: 0.9964 - val\_loss: 0.0742 - val\_accuracy: 0.9815

In [120]:

```
score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.07422380950862426

Test accuracy: 0.9815000295639038



In [121]:


```
w_after = model_batch.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10,10))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## ReLU + Adam Optimization + 2 Hidden Layers (with Dropout and Batch Normalization)

In [122]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras
```

```
from keras.layers import Dropout
```

```
model_drop = Sequential()
```

```
model_drop.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
```

```
model_drop.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
```

```
model_drop.add(Dense(output_dim, activation='softmax'))
```

```
model_drop.summary()
```

Model: "sequential\_22"

Layer (type)	Output Shape	Param #
dense_85 (Dense)	(None, 256)	200960
batch_normalization_33 (Batch Normalization)	(None, 256)	1024



dropout_11 (Dropout)	(None, 256)	0
dense_86 (Dense)	(None, 64)	16448
batch_normalization_34 (Batch Normalization)	(None, 64)	256
dropout_12 (Dropout)	(None, 64)	0
dense_87 (Dense)	(None, 10)	650

=====  
 Total params: 219,338  
 Trainable params: 218,698  
 Non-trainable params: 640  
 =====

In [123]:

```

model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 5s 76us/step - loss: 0.6283 - accuracy: 0.8107 - val_loss: 0.2089 - val_accuracy: 0.9392
Epoch 2/20
60000/60000 [=====] - 4s 71us/step - loss: 0.3279 - accuracy: 0.9043 - val_loss: 0.1595 - val_accuracy: 0.9493
Epoch 3/20
60000/60000 [=====] - 4s 70us/step - loss: 0.2703 - accuracy: 0.9201 - val_loss: 0.1373 - val_accuracy: 0.9561
Epoch 4/20
60000/60000 [=====] - 4s 70us/step - loss: 0.2330 - accuracy: 0.9319 - val_loss: 0.1144 - val_accuracy: 0.9645
Epoch 5/20
60000/60000 [=====] - 4s 70us/step - loss: 0.2040 - accuracy: 0.9405 - val_loss: 0.1108 - val_accuracy: 0.9649
Epoch 6/20
60000/60000 [=====] - 4s 70us/step - loss: 0.1898 - accuracy: 0.9439 - val_loss: 0.1002 - val_accuracy: 0.9708
Epoch 7/20
60000/60000 [=====] - 4s 71us/step - loss: 0.1766 - accuracy: 0.9481 - val_loss: 0.0982 - val_accuracy: 0.9711
Epoch 8/20
60000/60000 [=====] - 4s 71us/step - loss: 0.1669 - accuracy: 0.9509 - val_loss: 0.0899 - val_accuracy: 0.9727
Epoch 9/20
60000/60000 [=====] - 4s 70us/step - loss: 0.1548 - accuracy: 0.9549 - val_loss: 0.0859 - val_accuracy: 0.9743
Epoch 10/20
60000/60000 [=====] - 4s 69us/step - loss: 0.1464 - accuracy: 0.9570 - val_loss: 0.0898 - val_accuracy: 0.9731
Epoch 11/20
60000/60000 [=====] - 4s 69us/step - loss: 0.1405 - accuracy: 0.9590 - val_loss: 0.0900 - val_accuracy: 0.9749
Epoch 12/20
60000/60000 [=====] - 4s 69us/step - loss: 0.1362 - accuracy: 0.9589 - val_loss: 0.0889 - val_accuracy: 0.9746
Epoch 13/20
60000/60000 [=====] - 4s 70us/step - loss: 0.1289 - accuracy: 0.9618 - val_loss: 0.0836 - val_accuracy: 0.9749
Epoch 14/20
60000/60000 [=====] - 4s 72us/step - loss: 0.1225 - accuracy: 0.9635 - val_loss: 0.0785 - val_accuracy: 0.9776
Epoch 15/20
60000/60000 [=====] - 4s 70us/step - loss: 0.1216 - accuracy: 0.9641 - val_loss: 0.0795 - val_accuracy: 0.9771
Epoch 16/20
60000/60000 [=====] - 4s 70us/step - loss: 0.1141 - accuracy: 0.9664 - val_loss: 0.0785 - val_accuracy: 0.9785
Epoch 17/20
60000/60000 [=====] - 4s 70us/step - loss: 0.1106 - accuracy: 0.9674 - val_loss: 0.0753 - val_accuracy: 0.9779
Epoch 18/20

```

```
60000/60000 [=====] - 4s 71us/step - loss: 0.1100 - accuracy: 0.9667 - va
l_loss: 0.0762 - val_accuracy: 0.9786
Epoch 19/20
60000/60000 [=====] - 4s 71us/step - loss: 0.1027 - accuracy: 0.9696 - va
l_loss: 0.0739 - val_accuracy: 0.9789
Epoch 20/20
60000/60000 [=====] - 4s 70us/step - loss: 0.1020 - accuracy: 0.9692 - va
l_loss: 0.0713 - val_accuracy: 0.9795
```

In [124]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.0712568270204123  
Test accuracy: 0.9794999957084656



both training and validation loss keeps on decreasing with the increase in no of epochs and around 20epochs both the losses are pretty similar or almost the same

In [125]:

```
w_after = model_drop.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10,10))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
```

```
plt.title("Trained model Weights")
ax = sns.violinplot(y=hl_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## ReLU + Adam Optimization + 3 Hidden Layers (with Dropout and Batch Normalization)

In [126]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
```

```

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55,
seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()

```

Model: "sequential\_23"

Layer (type)	Output Shape	Param #
dense_88 (Dense)	(None, 512)	401920
batch_normalization_35 (Batch Normalization)	(None, 512)	2048
dropout_13 (Dropout)	(None, 512)	0
dense_89 (Dense)	(None, 256)	131328
batch_normalization_36 (Batch Normalization)	(None, 256)	1024
dropout_14 (Dropout)	(None, 256)	0
dense_90 (Dense)	(None, 128)	32896
batch_normalization_37 (Batch Normalization)	(None, 128)	512
dropout_15 (Dropout)	(None, 128)	0
dense_91 (Dense)	(None, 10)	1290
Total params: 571,018		
Trainable params: 569,226		
Non-trainable params: 1,792		

In [127]:

```

model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/20
60000/60000 [=====] - 10s 163us/step - loss: 0.7715 - accuracy: 0.7613 - val_loss: 0.2116 - val_accuracy: 0.9358
Epoch 2/20
60000/60000 [=====] - 9s 155us/step - loss: 0.3486 - accuracy: 0.8948 - val_loss: 0.1556 - val_accuracy: 0.9508
Epoch 3/20
60000/60000 [=====] - 10s 163us/step - loss: 0.2721 - accuracy: 0.9203 - val_loss: 0.1334 - val_accuracy: 0.9579
Epoch 4/20
60000/60000 [=====] - 10s 159us/step - loss: 0.2331 - accuracy: 0.9324 - val_loss: 0.1208 - val_accuracy: 0.9618
Epoch 5/20
60000/60000 [=====] - 9s 153us/step - loss: 0.2010 - accuracy: 0.9418 - val_loss: 0.1079 - val_accuracy: 0.9667
Epoch 6/20
60000/60000 [=====] - 9s 154us/step - loss: 0.1785 - accuracy: 0.9483 - val_loss: 0.0947 - val_accuracy: 0.9722
Epoch 7/20
60000/60000 [=====] - 9s 155us/step - loss: 0.1664 - accuracy: 0.9522 - val_loss: 0.0903 - val_accuracy: 0.9724
Epoch 8/20
60000/60000 [=====] - 9s 154us/step - loss: 0.1541 - accuracy: 0.9549 - val_loss: 0.0878 - val_accuracy: 0.9734
Epoch 9/20
60000/60000 [=====] - 9s 155us/step - loss: 0.1464 - accuracy: 0.9574 - val_loss: 0.0866 - val_accuracy: 0.9736

```

```

al_loss: 0.0868 - val_accuracy: 0.9733
Epoch 10/20
60000/60000 [=====] - 9s 155us/step - loss: 0.1354 - accuracy: 0.9606 - v
al_loss: 0.0787 - val_accuracy: 0.9751
Epoch 11/20
60000/60000 [=====] - 9s 156us/step - loss: 0.1247 - accuracy: 0.9627 - v
al_loss: 0.0794 - val_accuracy: 0.9764
Epoch 12/20
60000/60000 [=====] - 9s 153us/step - loss: 0.1199 - accuracy: 0.9639 - v
al_loss: 0.0752 - val_accuracy: 0.9782
Epoch 13/20
60000/60000 [=====] - 9s 156us/step - loss: 0.1174 - accuracy: 0.9660 - v
al_loss: 0.0777 - val_accuracy: 0.9780
Epoch 14/20
60000/60000 [=====] - 9s 154us/step - loss: 0.1095 - accuracy: 0.9674 - v
al_loss: 0.0719 - val_accuracy: 0.9781
Epoch 15/20
60000/60000 [=====] - 9s 156us/step - loss: 0.1065 - accuracy: 0.9680 - v
al_loss: 0.0706 - val_accuracy: 0.9784
Epoch 16/20
60000/60000 [=====] - 9s 154us/step - loss: 0.0971 - accuracy: 0.9708 - v
al_loss: 0.0696 - val_accuracy: 0.9807
Epoch 17/20
60000/60000 [=====] - 9s 154us/step - loss: 0.0960 - accuracy: 0.9717 - v
al_loss: 0.0663 - val_accuracy: 0.9814
Epoch 18/20
60000/60000 [=====] - 9s 155us/step - loss: 0.0948 - accuracy: 0.9719 - v
al_loss: 0.0713 - val_accuracy: 0.9797
Epoch 19/20
60000/60000 [=====] - 9s 155us/step - loss: 0.0900 - accuracy: 0.9731 - v
al_loss: 0.0641 - val_accuracy: 0.9831
Epoch 20/20
60000/60000 [=====] - 9s 154us/step - loss: 0.0869 - accuracy: 0.9742 - v
al_loss: 0.0657 - val_accuracy: 0.9818

```

In [128]:

```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.06569366658385843  
Test accuracy: 0.9818000197410583





In [129]:


```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10,10))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



# ReLU + Adam Optimization + 3 Hidden Layers (with Dropout and Batch Normalization)

In [130]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Model: "sequential\_24"

Layer (type)	Output Shape	Param #
=====		
dense_92 (Dense)	(None, 512)	401920
batch_normalization_38 (Batch Normalization)	(None, 512)	2048
dropout_16 (Dropout)	(None, 512)	0
dense_93 (Dense)	(None, 256)	131328
batch_normalization_39 (Batch Normalization)	(None, 256)	1024
dropout_17 (Dropout)	(None, 256)	0
dense_94 (Dense)	(None, 128)	32896
batch_normalization_40 (Batch Normalization)	(None, 128)	512
dropout_18 (Dropout)	(None, 128)	0
dense_95 (Dense)	(None, 64)	8256
batch_normalization_41 (Batch Normalization)	(None, 64)	256
dropout_19 (Dropout)	(None, 64)	0
dense_96 (Dense)	(None, 32)	2080
batch_normalization_42 (Batch Normalization)	(None, 32)	128
dropout_20 (Dropout)	(None, 32)	0

dense_97 (Dense)	(None, 10)	330
=====		
Total params: 580,778		
Trainable params: 578,794		
Non-trainable params: 1,984		
=====		

In [131]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 11s 186us/step - loss: 2.0119 - accuracy: 0.3304 -
val_loss: 0.7607 - val_accuracy: 0.8444
Epoch 2/20
60000/60000 [=====] - 10s 166us/step - loss: 1.0764 - accuracy: 0.6406 -
val_loss: 0.3660 - val_accuracy: 0.9143
Epoch 3/20
60000/60000 [=====] - 10s 166us/step - loss: 0.7278 - accuracy: 0.7714 -
val_loss: 0.2483 - val_accuracy: 0.9346
Epoch 4/20
60000/60000 [=====] - 10s 166us/step - loss: 0.5533 - accuracy: 0.8357 -
val_loss: 0.2000 - val_accuracy: 0.9451
Epoch 5/20
60000/60000 [=====] - 10s 165us/step - loss: 0.4630 - accuracy: 0.8702 -
val_loss: 0.1681 - val_accuracy: 0.9537
Epoch 6/20
60000/60000 [=====] - 10s 165us/step - loss: 0.3948 - accuracy: 0.8964 -
val_loss: 0.1581 - val_accuracy: 0.9583
Epoch 7/20
60000/60000 [=====] - 10s 166us/step - loss: 0.3486 - accuracy: 0.9102 -
val_loss: 0.1394 - val_accuracy: 0.9635
Epoch 8/20
60000/60000 [=====] - 10s 167us/step - loss: 0.3130 - accuracy: 0.9202 -
val_loss: 0.1312 - val_accuracy: 0.9672
Epoch 9/20
60000/60000 [=====] - 10s 170us/step - loss: 0.2941 - accuracy: 0.9265 -
val_loss: 0.1244 - val_accuracy: 0.9688
Epoch 10/20
60000/60000 [=====] - 10s 171us/step - loss: 0.2699 - accuracy: 0.9320 -
val_loss: 0.1168 - val_accuracy: 0.9719
Epoch 11/20
60000/60000 [=====] - 10s 167us/step - loss: 0.2502 - accuracy: 0.9392 -
val_loss: 0.1138 - val_accuracy: 0.9739
Epoch 12/20
60000/60000 [=====] - 10s 167us/step - loss: 0.2459 - accuracy: 0.9414 -
val_loss: 0.1119 - val_accuracy: 0.9746
Epoch 13/20
60000/60000 [=====] - 10s 166us/step - loss: 0.2271 - accuracy: 0.9462 -
val_loss: 0.1104 - val_accuracy: 0.9750
Epoch 14/20
60000/60000 [=====] - 10s 166us/step - loss: 0.2120 - accuracy: 0.9502 -
val_loss: 0.1035 - val_accuracy: 0.9737
Epoch 15/20
60000/60000 [=====] - 10s 166us/step - loss: 0.2085 - accuracy: 0.9510 -
val_loss: 0.1047 - val_accuracy: 0.9763
Epoch 16/20
60000/60000 [=====] - 10s 166us/step - loss: 0.1989 - accuracy: 0.9536 -
val_loss: 0.1012 - val_accuracy: 0.9760
Epoch 17/20
60000/60000 [=====] - 10s 168us/step - loss: 0.1892 - accuracy: 0.9559 -
val_loss: 0.1004 - val_accuracy: 0.9776
Epoch 18/20
60000/60000 [=====] - 10s 167us/step - loss: 0.1859 - accuracy: 0.9575 -
val_loss: 0.0933 - val_accuracy: 0.9787
Epoch 19/20
60000/60000 [=====] - 10s 166us/step - loss: 0.1772 - accuracy: 0.9597 -
val_loss: 0.0994 - val_accuracy: 0.9781
Epoch 20/20
60000/60000 [=====] - 10s 166us/step - loss: 0.1726 - accuracy: 0.9591 -
val_loss: 0.1020 - val_accuracy: 0.9760
```



```
val_loss: 0.1029 - val_accuracy: 0.9760
```

In [132]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

```
Test score: 0.10285850774878635
Test accuracy: 0.9760000109672546
```



In [133]:

```
w_after = model_drop.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10,10))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

#We have 5 hidden layers but we are only plotting distribution of only two layers
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
```

```
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## Comparing models using prettytable library

### Without using Batch Normalization and Dropout

In [140]:

```
from prettytable import PrettyTable
x=PrettyTable()

x.field_names=["#Layers", "Architecture", "Activation", "Optimizer", "Intializer", "Test
Score/Accuracy"] #column headers

x.add_row(["2", "256-64", "Relu", "Adam", "RandomNormal", "0.0899/0.980"])
x.add_row(["3", "512-256-128", "Relu", "Adam", "RandomNormal", "0.102/0.980"])
x.add_row(["5", "512-256-128-64-32", "Relu", "Adam", "RandomNormal", "0.076/0.982"])
print(x)
```

#Layers	Architecture	Activation	Optimizer	Intializer	Test Score/Accuracy
2	256-64	Relu	Adam	RandomNormal	0.0899/0.980
3	512-256-128	Relu	Adam	RandomNormal	0.102/0.980
5	512-256-128-64-32	Relu	Adam	RandomNormal	0.076/0.982

## Using Batch Normalization

In [139]:

```
from prettytable import PrettyTable
x=PrettyTable()

x.field_names=["#Layers", "Architecture", "Activation", "Optimizer", "Intializer", "Test
Score/Accuracy"] #column headers

x.add_row(["2", "256-64", "Relu", "Adam", "RandomNormal", "0.0934/0.976"])
x.add_row(["3", "512-256-128", "Relu", "Adam", "RandomNormal", "0.076/0.981"])
x.add_row(["5", "512-256-128-64-32", "Relu", "Adam", "RandomNormal", "0.074/0.981"])
print(x)
```

#Layers	Architecture	Activation	Optimizer	Intializer	Test Score/Accuracy
2	256-64	Relu	Adam	RandomNormal	0.0934/0.976
3	512-256-128	Relu	Adam	RandomNormal	0.076/0.981
5	512-256-128-64-32	Relu	Adam	RandomNormal	0.074/0.981

## With using Batch Normalization and Dropout

In [138]:

```
from prettytable import PrettyTable
x=PrettyTable()

x.field_names=["#Layers", "Architecture", "Activation", "Optimizer", "Intializer", "Test
Score/Accuracy"] #column headers

x.add_row(["2", "256-64", "Relu", "Adam", "RandomNormal", "0.0712/0.979"])
x.add_row(["3", "512-256-128", "Relu", "Adam", "RandomNormal", "0.065/0.981"])
x.add_row(["5", "512-256-128-64-32", "Relu", "Adam", "RandomNormal", "0.102/0.976"])
print(x)
```

#Layers	Architecture	Activation	Optimizer	Intializer	Test Score/Accuracy
2	256-64	Relu	Adam	RandomNormal	0.0712/0.979
3	512-256-128	Relu	Adam	RandomNormal	0.065/0.981
5	512-256-128-64-32	Relu	Adam	RandomNormal	0.102/0.976

In [ ]: